



---

# Lecture 19 – Optimization for GPUs

Jaejin Lee

Dept. of Computer Science and Engineering, College of Engineering

Dept. of Data Science, Graduate School of Data Science

Seoul National University

<http://aces.snu.ac.kr>



## References

- John Cheng, Max Grossman, and Ty McKercher, Professional CUDA C Programming, John Wiley & Sons, 2014




# Hardware Scheduling Units in GPUs

- Basic unit of GPUs for scheduling
  - All threads in it processes a single instruction at the same time in SIMD fashion
  - Lock-step
- NVIDIA - warp
  - 32 hardware threads (work-items)
- AMD - frontend
  - 64 hardware threads (work-items)
- Other vendors may have different names

# An SM in NVIDIA Fermi Architecture

- Each SM supports concurrent execution of hundreds of threads
- Multiple SMs per GPU
  - Thousands of threads executing concurrently on a single GPU
  - Multiple thread blocks may be assigned to the same SM at once
    - Scheduled based on the availability of SM resources





# SIMT

- NVIDIA terminology
- Single Instruction Multiple Thread (SIMT) architecture
  - To manage and execute threads in groups of 32 called warps
- All threads in a warp execute the same instruction at the same time
- Each SM partitions the assigned thread blocks into warps
  - The SM schedules the warps for execution on available hardware resources
  - All threads in a thread block run logically in parallel
    - Not all threads can execute physically at the same time
    - Different threads in a thread block may make progress at a different pace



# Thread Block Scheduling

- A thread block is scheduled on only one SM
  - The block remains there until execution completes
- An SM can hold more than one thread block at the same time
- Shared memory is partitioned among thread blocks resident on the SM
- Registers are partitioned among threads



## Thread Block Scheduling (cont'd)

- Sharing data among parallel threads may cause a race condition
  - Multiple threads accessing the same data with an undefined ordering, which results in unpredictable program behavior
  - CUDA provides a means to synchronize threads within a thread block
  - No primitives are provided for inter-block synchronization



# Warps and Thread Blocks

- Threads in the thread block are further partitioned into warps
  - A warp consists of 32 consecutive threads
  - All threads in a warp are executed SIMT fashion
  - All threads execute the same instruction
  - Each thread carries out that operation on its own private data
- Threads with consecutive values for threadIdx.x are grouped into warps
- For example, a one-dimensional thread block with 128 threads will be organized into 4 warps as follows:
  - Warp 0: thread 0, thread 1, thread 2, ... thread 31
  - Warp 1: thread 32, thread 33, thread 34, ... thread 63
  - Warp 3: thread 64, thread 65, thread 66, ... thread 95
  - Warp 4: thread 96, thread 97, thread 98, ... thread 127



## Warps and Thread Blocks (cont'd)

- The logical layout of a two or three-dimensional thread block can be converted into its one-dimensional physical layout
  - Using the  $x$  dimension as the innermost dimension, the  $y$  dimension as the second dimension, and the  $z$  dimension as the outermost
- For example, given a 2D thread block, a unique identifier for each thread in a block can be calculated using the built-in **threadIdx** and **blockDim** variables:
  - **threadIdx.y \* blockDim.x + threadIdx.x**
- The same calculation for a 3D thread block is as follows:
  - **threadIdx.z \* blockDim.y \* blockDim.x**
  - **+ threadIdx.y \* blockDim.x + threadIdx.x**




## Warp Scheduling

- Warps within a thread block may be scheduled in any order
- The number of active warps is limited by SM resources
- When a warp idles for any reason (e.g., waiting for values to be read from device memory), the SM is free to schedule another available warp from any thread block that is resident on the same SM
- Switching between concurrent warps has no overhead
  - Hardware resources are partitioned among all threads and blocks on an SM


## Warp Scheduling (cont'd)

- Each SM features two warp schedulers and two instruction dispatch units
- When a thread block is assigned to an SM, all threads in a thread block are divided into warps
  - The two warp schedulers select two warps and issue one instruction from each warp to a group of 16 CUDA cores, 16 load/store units, or four special function units
  - The Fermi architecture, compute capability 2. $x$ , can simultaneously handle 48 warps per SM for a total of 1,536 threads resident in a single SM at a time



# Branch Divergence

- GPUs do not have complex branch prediction mechanisms
- All threads in a warp must execute identical instructions on the same cycle
  - If one thread executes an instruction, all threads in the warp must execute that instruction
  - This could become a problem if threads in the same warp take different paths through an application
    - When work-items in the same work-group follow different paths of control flow, they diverge in their execution
    - If - then - else
    - Loops with different loop bounds for different threads
  - A low degree of divergence will be better
  - Pick a work-group size that is a multiple of the warp size





## Resources for Warps

- The local execution context of a warp mainly consists of the following resources:
  - Program counters
  - Registers
    - Each SM has a set of 32-bit registers stored in a register file that are partitioned among threads
    - A few thousands of registers
  - Shared memory (local memory in OpenCL)
    - A fixed amount of shared memory that is partitioned among thread blocks
- The execution context of each warp processed by an SM is maintained on-chip during the entire lifetime of the warp
  - Switching from one execution context to another has no cost



## Resources for Warps (cont'd)

- The number of thread blocks and warps that can simultaneously reside on an SM for a given kernel depends on the number of registers and amount of shared memory available on the SM and required by the kernel
  - When each thread consumes more registers, fewer warps can be placed on an SM
  - When a thread block consumes more shared memory, fewer thread blocks are processed simultaneously by an SM



## Thread Granularity

- Put more work into each thread (work-item) and use fewer threads (work-items)
  - May reduce the kernel launching overhead
  - May remove redundant computations between threads (work-items)
  - May increase the number of registers, resulting in low occupancy
  - May reduce the number of thread blocks (work-groups), resulting in making the SM underutilized



## Grid and Block Size

- Keep the number of threads per block a multiple of warp size (32)
- Avoid small block sizes
  - Start with at least 128 or 256 threads per block
- Adjust block size up or down according to kernel resource requirements
- Keep the number of blocks much greater than the number of SMs to expose sufficient parallelism
- Conduct experiments to discover the best execution configuration and resource usage





# Synchronization

- Two levels:
  - System level: wait for all work on both the host and the device to complete
    - `cudaDeviceSynchronize()` can be used to block the host until all CUDA operations have been completed
  - Block level: wait for all threads in a thread block to reach the same point in execution
    - `__syncthreads()`

## Parallel Reduction


- Partition the input vector into smaller chunks
- Have a thread calculate the partial sum for each chunk
- Add the partial results from each chunk into a final sum

```
int sum = 0;  
for (int i = 0; i < N; i++)  
    sum += array[i];
```



# Occupancy

- The number of active warps per Streaming Multiprocessor (AMD calls it a Streaming Core)
  - Computed at compile time
- It describes how well the resources of the SM are being utilized





## Occupancy (cont'd)

- A thread block is called an active block when compute resources, such as registers and shared memory, have been allocated to it
  - The warps it contains are called active warps
  - The warp schedulers on an SM select active warps on every cycle and dispatch them to execution units
- Active warps can be further classified into the following three types:
  - Selected warp: an active warp that is actively executing
  - Stalled warp: an active warp that is ready for execution but not currently executing
  - Eligible warp: an active warp that is not ready for execution



## Occupancy (cont'd)

- A warp is eligible for execution if both of the following two conditions are met:
  - Thirty-two CUDA cores are available for execution
  - All arguments to the current instruction are ready
- You need to keep a large number of warps active to hide the latency caused by warps stalling
- Occupancy is the ratio of active warps to maximum number of warps, per SM
  - $$\text{occupancy} = \frac{\# \text{ of active warps}}{\# \text{ of maximum warps}}$$
  - Less than or equal to 1, the bigger, the better



## Occupancy (cont'd)

- To enhance occupancy,
  - Resize the thread block configuration or re-adjust resource usage to permit more simultaneously active warps
  - Small thread blocks: too few threads per block leads to hardware limits on the number of warps per SM to be reached before all resources are fully utilized
  - Large thread blocks: too many threads per block leads to fewer per-SM hardware resources available to each thread



# Global Memory Accesses

- Global memory loads/stores are cached
  - All accesses to global memory go through the L2 cache
  - Many accesses also pass through the L1 cache, depending on the type of access and the GPU's architecture
- Kernel memory requests are typically served between the device DRAM and SM on-chip memory using either 128-byte or 32-byte memory transactions
  - If both L1 and L2 caches are used, a memory access is serviced by a 128-byte memory transaction
  - If only the L2 cache is used, a memory access is serviced by a 32-byte memory transaction



## Global Memory Accesses (cont'd)

- On architectures that allow the L1 cache to be used for global memory caching, the L1 cache can be explicitly enabled or disabled at compile time
- An L1 cache line is 128 bytes, and it maps to a 128-byte aligned segment in device memory
  - If each thread in a warp requests one 4-byte value, that results in 128 bytes of data per request, which maps perfectly to the cache line size and device memory segment size




# Aligned and Coalesced Memory Accesses

- Aligned memory accesses
  - When the first address of a device memory transaction is an even multiple of the cache granularity being used to service the transaction
  - Either 32 bytes for L2 cache or 128 bytes for L1 cache
  - Performing a misaligned load will cause wasted bandwidth
- Coalesced memory accesses
  - When all 32 threads in a warp access a contiguous chunk of memory


## Aligned and Coalesced Memory Accesses (cont'd)

- To maximize global memory throughput, aligned coalesced memory accesses are ideal
  - A warp accessing a contiguous chunk of memory starting at an aligned memory address
- For example, aligned and coalesced memory load operations require only a single 128-byte memory transaction to read the data from device memory



## Aligned and Coalesced Memory Accesses (cont'd)

- For example, misaligned and uncoalesced memory accesses require as many as three 128-byte memory transactions to read the data from device memory
  - One starting at offset 0
  - One starting at offset 128
  - One starting at offset 256






## Host-device Data Transfer

- Data transfer between the host and the device has a much lower bandwidth than global memory access
  - PCI-E: a few GB/s
  - Global memory: a few hundred GB/s
- Minimize data transfer between the host and the device
  - Better to recompute on the accelerator
  - Use the global memory on the accelerator for intermediate data
- One large transfer is much faster than many small ones

# Asynchronous Copy

- Overlap of data transfer and code execution
  - Simultaneously execute a kernel while performing a copy between the device and host memory
- Relevant only for accelerators using PCI-E bus
- Use two queues (two CUDA streams)
  - One for data transfer and one for execution
  - Use events to enforce dependences



# Pipelining (Double Buffering)

- Vector addition example
  - $A + B = C$
  - Divide large vectors into segments
    - $A_1, A_2, \dots, A_n$
    - $B_1, B_2, \dots, B_n$
    - $C_1, C_2, \dots, C_n$
- Overlap transfer and addition

